

Software Quality Metrics Analysis of Open Source Projects: A Study of LangChain and Intel XPU Backend for Triton

Garrett Gruss

Southern Methodist University

CS8314-795B-1257

December 5, 2025

Abstract

This report presents a study of software quality metrics pertaining to two popular open source projects: `Langchain-ai/langchain` and `intel/intel-xpu-backend-for-triton`. The initial goal of this study was to collect discrete internal quality metrics such as lines of code and cyclomatic complexity over time, but the project pivoted to focus on DORA (DevOps Research and Assessment) metrics due to permissions issues. Pull requests, issues, and merge patterns over a timeframe of one month were collected from GitHub project insights. These metrics were used to calculate deployment frequency, lead time for changes, change failure rate and time to restore service. This study demonstrates the value of using DORA metrics as practical indicators of software project health and team performance in open source development environments.

1 Introduction

1.1 Background

The automated measurement of software quality metrics is essential for proactively monitoring and improving development processes. As open source projects grow in complexity and importance, the ability to identify and quantify software quality becomes increasingly important for maintainers, contributors, and organizations depending on these projects.

The importance of routine, repeatable quality measurement can be extended to frameworks such as Goal-Question-Metrics (GQM) and the Quality Improvement Paradigm (QIP). These frameworks formalize the process of translating organizational goals into measurable indicators, enabling corrective action and discrete quality tracking. In modern DevOps environments, team velocity can necessitate multiple deployments per day, issue resolutions, and defect identifications. This rapid velocity and continuous integration and deployment can be modeled and tracked using metrics like the DORA framework's deployment frequency, lead time for changes, change failure rate, and time to restore service have become industry standards for assessing team performance and delivery capability. For open source projects composed of distributed teams collaborating across the globe, the automated collection of quality metrics becomes essential to manage the project's health, productivity, and code quality while maintaining the rapid velocity that bleeding-edge projects require.

1.2 Motivation

The motivation for this project was the need to systematically assess the software quality of popular open source projects. In analyzing established projects with highly active communities, we can identify best practices in the open source community, identify patterns that correlate with project success, provide insights into development velocity, maintainability, and reliability, and establish reusable benchmarks for software quality.

1.3 Project Objectives

The primary objectives of this project were to collect software quality metrics from two major open source projects over a timespan of one month, analyze these metrics to assess project health, compare these metrics across projects to identify trends and patterns, and provide actionable insights based on quantitative analysis.

1.4 Selected Projects

1.4.1 LangChain ([langchain-ai/langchain](https://github.com/langchain-ai/langchain))

LangChain is a highly popular framework for developing durable and feature rich applications powered by large language models (LLMs). The project is built in Python, has 128k stars, 20k forks, and 3,812 active contributors.

1.4.2 Intel XPU Backend for Triton ([intel/intel-xpu-backend-for-triton](https://github.com/intel/intel-xpu-backend-for-triton))

Intel XPU Backend for Triton is an out-of-tree backend module for Triton designed to provide high-performance GPU computing on Intel GPUs for PyTorch and standalone usage. The project targets Intel Data Center GPU Max Series, Flex Series, and Arc GPUs, supporting Ubuntu operating systems. As a compiler backend, it uses LLVM to generate optimized code for Intel GPUs, enabling developers to write high-performance kernels using Triton's Python-based domain-specific language. The project is actively maintained as part of Intel's deep learning ecosystem and provides an alternative backend for the Triton framework. This project has 222 stars, 78 forks, and 538 active contributors.

2 Related Work

The measurement of software quality metrics in continuous integration and deployment environments has been studied from both internal and external metric perspectives. This section

identifies relevant research that informs the methodology and metrics selection criteria used within this paper.

2.1 Internal Metrics and Static Analysis

Zampetti et al. [1] investigated how open source projects utilize static code analysis tools such as Checkstyle, PMD, and FindBugs into their pipelines. Zampetti et al. found 97% of build warnings and 6% of build failures were triggered by style issues, opposed to defects such as test failures. This study described how static code analysis in CI/CD pipelines can be utilized to identify defects and warnings within the code base prior to deployment.

[2] developed the Open Code Coverage Framework (OCCF) to formalize coverage measurement across different programming languages. Their framework reduced implementation requirements by roughly 90% while supporting multiple coverage types such as statement, decision, branch, and path. This study provided a framework for software quality metrics across a host of different programming languages.

Yu et al. [3] examined the visualization of metrics collected through the CI/CD process. They interviewed 22 participants with visualizations of 14 base metrics. Their key finding was 86% of participants appreciated the visualization of dashboard CI-data, but found a majority of participants complained about information overload, highlighting the need to focus on key metrics rather than all obtainable metrics.

2.2 External Metrics and DevOps Performance

The DevOps Research and Assessment (DORA) metrics framework has emerged as an industry-standard indicator of cloud native software delivery performance. Wilkes et al. [4] presented a framework for automatically extracting DORA metrics such as Deployment Frequency, Lead Time for Changes, Mean Time to Recover, Change Failure Rate. Their analysis of 304 GitHub projects established performance baselines: top-quartile projects achieve 116.2 deployments per year with 5.1-day lead times, while bottom-quartile projects show 6.6 de-

ploys per year with 140.6-day lead times. These metrics provide a useful baseline for the DORA performance of any future project under analysis.

2.3 Security Measurement

Zhuravchak et al. [5] addressed temporal security measurement by utilizing a persistent visualization of security trends across project builds. Their approach combined Trivy scanning, GitHub actions automation, and Chart.js to visualize how dependency updates and code modification influence the security posture of the system over time.

2.4 Traceability Automated Metrics Collection

Stahl et al. [6] documented the Eiffel framework for standardized traceability throughout CI/CD pipelines. The Eiffel framework uses a JSON-based schema to organize event messages into a directed acyclic graph (DAG). This approach produces traceable events that can be standardized across the organization to provide deep insights into the performance of distributed projects. Unfortunately, the Eiffel framework is no longer active, and has been replaced with the OpenTelemetry community.

OpenTelemetry [7] provides vendor-neutral logging, tracing, and metrics collection through the OpenTelemetry Logging Protocol (OTLP). The OpenTelemetry GitHub receiver [8] enables automated scrapping of version control system (VCS) metrics via GitHub’s API, enabling continuous automated measurement of internal code quality metrics without manual data extraction. Kalliamvakou et al. [9] demonstrated the methodology and feasibility of extracting development process metrics from VCS, though they identified challenges in data completeness. This project leveraged the OpenTelemetry GitHub receiver to collect VCS metrics aligned with the OpenTelemetry semantic conventions v1.37.0.

3 Methodology

This project evolved through three distinct phases, each representing an operational pivot in the metric collection methodology in response to technical and time-based constraints. This section documents the evolution of this study from initial planning through final implementation, providing insight into future decision-making during the design and execution of any metrics collection activities.

3.1 Phase 1: Original Plan - Static Code Analysis via GitHub Actions

3.1.1 Initial Objectives

The original methodology aimed to collect internal quality metrics within the code by integrating custom GitHub Actions workflows into target repositories. These workflows would trigger static code analysis tools on each commit or pull request to extract time-based internal metrics. The metrics chosen for this analysis were Lines Of Code (LOC), representing the source, comment, and blank lines measured via tools like cloc or tokei. Cyclomatic Complexity would be measured code path complexity measured via radon for Python. Code Coverage would be measured via a tool like pytest-cov, computed from the percentage of code exercised by tests. Technical debt would be estimated by remediation time, as calculated from SonarQube or CodeClimate analysis. Maintainability index would be computed as a composite metrics combining complexity, LOC, and comment density.

3.1.2 Planned Implementation

The planned implementation of the static collection of metrics via workflows involved creating custom GitHub Actions workflows, installing the custom static analysis tools in the workflow environment, running the analysis on each commit or pull request, and publishing the metrics to a centralized database or observability platform.

3.1.3 Abandonment Rationale

This approach was abandoned during implementation due to its invasive nature. Adding custom GitHub actions and workflows to projects required forking repositories on a per-branch basis, resulting in excessive modification of the source code repository to collect metrics. Large projects like Langchain also have existing CI/CD pipelines, so integrating custom Actions involved risking compatibility conflicts with existing workflows, which could pollute the accuracy of collected software quality metrics. Each target repository would also required custom workflows tailored to it's language, build system, and testing framework. The invasiveness of this approach required a pivot to a more observant solution.

3.2 Phase 2: First Pivot - DORA Metrics via OpenTelemetry

3.2.1 Revised Approach

To avoid modification to the target repository, the collection methodology was pivoted to utilizing the Open Telemetry collectors to scrape DORA metrics from GitHub. This non-invasive approach automated metric collection through the use of GitHub's GraphQL and REST APIs without requiring change to the target repositories. The planned system was deployed through docker Compose using an OpenTelemetry Collector container that scraped GitHub repositories every 60 seconds and an OpenObserve metrics aggregation and visualization service that collected the streamed metrics.




Figure 1: OpenTelemetry DORA Metrics Architecture (Phase 2 Approach)

3.2.2 Target Metrics

The GitHub receiver was configured to collect nine VCS (Version Control System) metrics conforming to OpenTelemetry semantic conventions v1.37.0:

Deployment Frequency Support:

- `vcs.change.count` - Pull request counts by state (open/merged)
- `vcs.repository.count` - Total repositories monitored

Lead Time for Changes Support:

- `vcs.change.time_to_merge` - Duration from PR creation to merge (seconds)

- `vcs.change.time_to_approval` - Duration from PR creation to approval (seconds)
- `vcs.change.duration` - Time PRs remain in open state (seconds)

Development Velocity Indicators:

- `vcs.ref.count` - Branch/tag counts per repository
- `vcs.ref.time` - Branch lifespan from creation (seconds)
- `vcs.ref.lines_delta` - Code churn (lines added/removed) vs. trunk
- `vcs.ref.revisions_delta` - Commit count ahead/behind trunk

3.2.3 Challenges Encountered

While less invasive than GitHub Actions, the OpenTelemetry approach encountered permissions issues when scraping metrics from public repositories. The OpenTelemetry collector requires a GitHub PAT token with specific organization access to **repo** and **read:org**. Public repositories limit provided data through unauthorized API access, prohibiting the collection of detailed PR metrics required to compute DORA. Accessing langchain-ai and intel repositories requires a GitHub PAT with organization access, approved by organization administrators.

3.3 Phase 3: Final Pivot - Manual DORA Metrics from GitHub Web UI

3.3.1 Rationale for Final Approach

Since automated collection of detailed metrics was prohibited due to permissions issues, a simpler approach of scraping metrics from the GitHub Insights page was chosen as the preferred data collection strategy. This approach leveraged publicly accessible GitHub repository insights through the web interface. While less automated and less detailed than the

scraped metrics, this approach did not require specific permissions granted by organizational administrators.

3.3.2 Data Collection Process

Data was manually extracted from GitHub’s web interface using the repository Insights tab and Pull Requests/Issues pages. For each target repository, Pull Request Events, Issue Events, and metadata were collected. This manual approach bypassed all authentication and permission barriers encountered in previous phases.

3.3.3 Data Structure

The collected data was organized into TXT files:

- `opened_requests.txt` - Timestamps of PR creation
- `merged_requests.txt` - Timestamps of PR merges
- `closed_requests.txt` - Timestamps of PR closures
- `opened_issues.txt` - Timestamps of issue creation

3.4 Analysis Methodology

3.4.1 DORA Metrics Framework

The four key DORA metrics, established by the DevOps Research and Assessment team, serve as the analytical foundation:

1. **Deployment Frequency (DF):** How often an organization successfully releases to production (proxied by merged PR rate)
2. **Lead Time for Changes (LT):** The time it takes for a commit to get into production (measured as PR open-to-merge duration)

3. **Change Failure Rate (CFR):** The percentage of deployments causing a failure in production (calculated from bug issue frequency)
4. **Time to Restore Service (MTTR):** How long it takes to recover from a failure in production (measured as bug report to fix merge duration)

3.4.2 Metric Calculation

A Python-based metrics collection pipeline was constructed that loads and preprocesses scrapped metrics data, converts them into columnar CSV data, calculates DORA metrics using industry-standard formulas, performs statistical analysis, classifies DORA performance, and generates visualizations for each metric.

3.4.3 Deployment Frequency Calculation

$$DF = \frac{\text{Total Merged PRs}}{\text{Time Period (days)}} \quad (1)$$

3.4.4 Lead Time Calculation

For each merged PR i :

$$LT_i = T_{\text{merged}_i} - T_{\text{opened}_i} \quad (2)$$

Aggregate metric: $\text{median}(LT_i)$

3.4.5 Change Failure Rate Calculation

$$CFR = \frac{\text{Number of Bug Issues}}{\text{Total Deployments}} \times 100\% \quad (3)$$

Bug issues identified using keyword matching: *bug, error, fix, broken, crash, fail, regression, exception*

3.4.6 Time to Restore Calculation

For each bug-fix pair (b, f) :

$$MTTR_{b,f} = T_{\text{fix merged}_f} - T_{\text{bug opened}_b} \quad (4)$$

Aggregate metric: $\text{median}(MTTR_{b,f})$

3.4.7 Performance Classification

DORA performance was classified according to the 2023 DORA State of DevOps Report standards.

Table 1: DORA Performance Level Classifications

Metric	Elite	High	Medium	Low
Deployment Frequency	Multiple/day	Weekly	Monthly	< Monthly
Lead Time	< 1 day	1-7 days	1-4 weeks	> 1 month
Change Failure Rate	0-15%	16-30%	31-45%	> 45%
Time to Restore	< 1 hour	1-7 days	1-7 days	> 1 week

4 Results

4.1 LangChain Analysis

4.1.1 Data Overview

Analysis of LangChain-ai/Langchain was conducted over a 30-day period from October 29, 2025 to November 27, 2025. During this period, 179 merged pull requests, 67 opened pull requests, 89 closed pull requests, and 69 issues were captured and analyzed.

4.1.2 Deployment Frequency

The LangChain project demonstrated high rate of deployment activity during the analysis period. The mean deployment frequency was 5.97, with a median of 8.5 deployments per day, resulting in 41.77 deployments per week on average. This deployment frequency classifies the project as an elite performer according to the DORA benchmark. This elite deployment frequency represents a mature CI/CD pipeline and team mentality resulting in multiple deploys per day.




Figure 2: LangChain Deployment Frequency Over Time

4.1.3 Lead Time for Changes

lead time for changes was estimated using temporal proximity matching between opened and merged PRs. The analysis across 162 matched PR pairs revealed a median lead time of 6.00 days and a mean of 4.98 days. The distribution showed a 25th percentile of 2.00 days and a 90th percentile of 7.00 days. Five pull requests were merged within 24 hours, representing

the quickest turnaround times. According to DORA performance standards, the median lead time of 6.00 days falls within the high performance category of 1-7 days.

4.1.4 Change Failure Rate

Change failure rate was estimated by using keyword matching on failure terms such as "bug", "error", "fix", "broken", "crash", "fail", "regression", and "exception". Out of 69 collected issues, 16 were classified as bug-related, representing 23.2% of the total population. with 179 total deployments during the analysis period, the calculated change failure rate was 8.94%. This rate falls within the elite performance category according to the DORA standard, which classifies rates between 0-15% as elite.




Figure 3: LangChain Deployments vs Bug Issues

4.1.5 Time to Restore Service

The time to restore service (MTTR) metric could not be calculated due to limitations occurring collecting related issues and merged pull requests. Estimation of MTTR requires collecting an opened issue, and identifying the time when the issue was resolved through a pull request and merge approval event. Future work could improve bug-fix pair identification using a properly authenticated GitHub PAT connected to an open-Telemetry GitHub scraper.

4.1.6 Overall Performance

Table 2: LangChain DORA Metrics Summary

Metric	Value	Classification
Deployment Frequency	5.97/day, 41.77/week	Elite
Lead Time for Changes	6.00 days (median)	High
Change Failure Rate	8.94%	Elite
Time to Restore Service	N/A	N/A
Overall Rating	High Performer	

4.2 Intel XPU Backend for Triton Analysis

4.2.1 Data Overview

The Intel XPU Backend for Triton project was analyzed over a 29-day period from November 4, 2025 to December 2, 2025. The dataset included 104 merged pull requests, 17 opened pull requests, 73 closed pull requests, and 42 opened issues.

4.2.2 Deployment Frequency

The mean deployment frequency was 3.59 deployments per day with a median of 9.0 deployments per day, totaling to 25.10 average deployments per week. The maximum number of deployments in a single day was 28. According to the DORA standard, this frequency classifies as an elite performer, achieving multiple deployments per day.

4.2.3 Lead Time for Changes

lead time for changes was estimated using temporal proximity matching between opened and merged PRs. However, the metrics collected were statistically insufficient, with only 3

matching pull requests identified. The limited sample size makes this metric unreliable for overall DORA assessment.

4.2.4 Change Failure Rate

Bug-related issues were identified using keyword matching. Of the total 42 issues collected, 11 were classified as defect-related, representing 26.2% of the total issue population. With 104 total deployments during the analysis period, the calculated failure rate was 10.58%. According to the DORA standard, this rate falls within the elite performance category of 0-15%.

4.2.5 Time to Restore Service

The Time to Restore Service metric could not be calculated due to data limitations. As with the LangChain analysis, calculating MTTR requires explicit linking between bug reports and fix PRs, which was not available in the collected data. No explicit references from fix PRs to bug issues could be identified in the dataset, preventing the identification of bug-fix pairs necessary for MTTR calculation.

4.2.6 Overall Performance

Table 3: Intel XPU Backend DORA Metrics Summary

Metric	Value	Classification
Deployment Frequency	3.59/day, 25.10/week	Elite
Lead Time for Changes	N/A (insufficient data)	N/A
Change Failure Rate	10.58%	Elite
Time to Restore Service	N/A	N/A
Overall Rating	Elite Performer*	

* Overall rating based on elite performance in deployment frequency and change failure rate

4.3 Comparative Analysis

Table 4: Comparative DORA Metrics: LangChain vs Intel XPU Backend

Metric	LangChain	Intel XPU
Deployment Frequency	5.97/day (Elite)	3.59/day (Elite)
Lead Time for Changes	6.00 days (High)	N/A
Change Failure Rate	8.94% (Elite)	10.58% (Elite)
Time to Restore Service	N/A	N/A
Overall Rating	High Performer	Elite Performer

5 Discussion

5.1 Key Findings

5.1.1 Development Velocity

Both projects achieve elite-level deployment frequency with LangChain performing at 5.97 deployments per day, and Intel XPU Backend at 3.59 daily deployments. These high rates indicate mature, highly active projects with robust CI/CD pipelines supporting multiple daily releases.

5.1.2 Code Quality and Reliability

Both projects achieve elite-level change failure rates below the 15% threshold: LangChain at 8.94%, Intel XPU Backend at 10.58%. High deployment frequency combined with low failure rates indicates robust testing, effective code review, and comprehensive CI/CD quality gates, demonstrating mature velocity-quality balance.

5.1.3 Incident Response

Time to restore service could not be calculated due to limitations in estimating which issue corresponded with a closed pull request.

5.2 Project Comparison

5.2.1 Similarities

Both projects achieved elite-level deployment frequency and change failure rates during daily development and merge activity. This pattern describes open source best practices regardless of domain or size.

5.2.2 Differences

Both projects shared an elite DORA classification, despite differences in popularity. Notably, the Intel project achieves 60% of Langchain’s deployment frequency despite only having 0.17% of the stars and 14% of contributors, suggesting higher productivity in Intel’s smaller corporate backed team.

5.3 Industry Benchmarking

Both projects significantly exceed Wilkes et al.’s top-quartile benchmark of 116.2 annual deployments (0.32 per day): LangChain with 2,180 annual deployments and Intel XPU Backend with 1,310, placing both in the top percentile. Elite change failure rates (under 15%) match DORA 2023 elite performer classification.

5.4 Practical Value of DORA Metrics

DORA metrics enable rapid comparison across projects through simple calculation and standardized classification. Unlike internal code metrics requiring language-specific tooling,

DORA metrics use version control and project management information. This enabled a simple and direct comparison between LangChain and Intel XPU Backend despite differences in domain, size, and architecture. Standardized performance tiers (elite/high/medium/low) enable immediate project assessment and visualization, without requiring a complex explanation of the intent and meaning behind the metrics. Practical applications include portfolio analysis, competitive benchmarking, and trend monitoring.

6 Lessons Learned

Deep analysis of software quality faces a hurdle of complexity and difficulty in collecting and aggregating metrics. GitHub Actions provided a rich avenue for code-level metrics, but required invasive repository modifications. Open Telemetry eliminated this invasiveness, but required repository permissions. Manual collection from publicly GitHub data proved sufficient for collection of DORA metrics, but is not practical for long running analysis of large projects.

7 Conclusion

This study focused on the capture and analysis of software quality metrics from two open source projects using DORA metrics from public GitHub data. Both LangChain and Intel XPU Backend achieved elite-level performance, despite their difference in popularity. DORA metrics provide a simple and standardized metric for direct comparison across projects regardless of domain, size, or architecture. This study’s comparison between an AI framework and compiler infrastructure demonstrates practical application for project benchmarking, comparison, and trend analysis.

Future work should focus on correlating DORA trends with internal metrics such as cyclomatic complexity, test coverage, and technical debt to establish predictive relationships between code quality and production quality. DORA metrics can be utilized as leading

indicators of quality degradation, and cross-domain comparison with internal metrics can be used to trace defects within the system to degradation of system behavior.

References

- [1] F. Zampetti, S. Scalabrino, R. Oliveto, G. Canfora, and M. Di Penta, “How open source projects use static code analysis tools in continuous integration pipelines,” in *2017 IEEE/ACM 14th International Conference on Mining Software Repositories (MSR)*. Buenos Aires, Argentina: IEEE, 2017, pp. 334–344.
- [2] K. Sakamoto, H. Washizaki, and Y. Fukazawa, “Open code coverage framework: A consistent and flexible framework for measuring test coverage supporting multiple programming languages,” in *2010 10th International Conference on Quality Software*. Zhangjiajie, China: IEEE, 2010, pp. 262–269.
- [3] L. Yu, E. Alégroth, P. Chatzipetrou, and T. Gorschek, “A roadmap for using continuous integration environments,” *Communications of the ACM*, vol. 67, no. 6, pp. 82–90, June 2024.
- [4] B. Wilkes, A. M. P. Milani, and M.-A. Storey, “A framework for automating the measurement of devops research and assessment (dora) metrics,” in *2023 IEEE International Conference on Software Maintenance and Evolution (ICSME)*. Bogotá, Colombia: IEEE, 2023, pp. 62–72.
- [5] Y. Zhuravchak, D. Zhuravchak, A. Zhuravchak, and I. Opirsky, “Security regression visualization in CI/CD using Trivy and GitHub Actions,” in *CSDP’2025: Cyber Security and Data Protection*, Lviv, Ukraine, July 2025, pp. 272–281.
- [6] D. Ståhl, K. Hallén, and J. Bosch, “Continuous integration and delivery traceability in industry: Needs and practices,” in *2016 42nd Euromicro Conference on Software*

Engineering and Advanced Applications (SEAA). Limassol, Cyprus: IEEE, 2016, pp. 68–72.

- [7] OpenTelemetry Authors, “OpenTelemetry specification,” Cloud Native Computing Foundation, 2021, accessed: 2025-04-19. [Online]. Available: <https://opentelemetry.io/docs/specs/otel/>
- [8] OpenTelemetry Collector Contributors, “GitHub receiver,” <https://github.com/open-telemetry/opentelemetry-collector-contrib/blob/main/receiver/githubreceiver/README.md>, 2024, openTelemetry Collector Contrib Repository. Accessed: 2025-04-19.
- [9] E. Kalliamvakou, G. Gousios, K. Blincoe, L. Singer, D. M. German, and D. Damian, “The promises and perils of mining GitHub,” in *Proceedings of the 11th Working Conference on Mining Software Repositories*, ser. MSR 2014. Hyderabad, India: ACM, 2014, pp. 92–101.

A Phase 1 Planned Configuration (GitHub Actions)

This section documents the original Phase 1 approach that involved integrating GitHub Actions workflows into target repositories for automated static code analysis and quality metric collection.

A.1 Planned GitHub Actions Workflow

The intended workflow would have been added to target repositories at `.github/workflows/quality-metr`

Listing 1: Planned quality-metrics.yml Workflow

```
1 name: Software Quality Metrics Collection
2
3 on:
```

```

4   push:
5       branches: [main, develop]
6   pull_request:
7       branches: [main]
8
9   jobs:
10      quality-analysis:
11          runs-on: ubuntu-latest
12
13          steps:
14              - name: Checkout code
15                uses: actions/checkout@v3
16
17              - name: Set up Python
18                uses: actions/setup-python@v4
19                with:
20                    python-version: '3.11'
21
22              - name: Install dependencies
23                run: |
24                    pip install radon pytest pytest-cov ruff mypy
25                    pip install -r requirements.txt
26
27              - name: Run Radon - LOC Analysis
28                run: |
29                    radon raw . -j > metrics/radon_loc.json
30                    radon cc . -j > metrics/radon_complexity.json
31                    radon mi . -j > metrics/radon_maintainability.json
32

```

```
33 - name: Run Pytest with Coverage
34   run: |
35     pytest --cov=. --cov-report=json \
36           --cov-report=term \
37           --junitxml=metrics/pytest_results.xml
38
39 - name: Run Ruff Linter
40   run: |
41     ruff check . --output-format=json \
42           > metrics/ruff_violations.json
43
44 - name: Run Mypy Type Checker
45   run: |
46     mypy . --json-report metrics/mypy_report
47
48 - name: Enforce Quality Gates
49   run: |
50     python .github/scripts/enforce_quality_gates.py
51
52 - name: Publish Metrics
53   uses: actions/upload-artifact@v3
54   with:
55     name: quality-metrics
56     path: metrics/
57
58 - name: Send to Observability Platform
59   run: |
60     python .github/scripts/publish_to_otel.py
61   env:
```

```
62     OTEL_ENDPOINT: ${ secrets.OTEL_ENDPOINT }}
63     OTEL_API_KEY: ${ secrets.OTEL_API_KEY }}
```

A.2 Planned Quality Gate Enforcement

The workflow would have included automated quality gates to prevent merging code that failed thresholds:

Listing 2: enforce_quality_gates.py

```
1  import json
2  import sys
3
4  # Load metrics
5  with open('metrics/radon_complexity.json') as f:
6      complexity = json.load(f)
7
8  with open('metrics/pytest_coverage.json') as f:
9      coverage = json.load(f)
10
11 with open('metrics/ruff_violations.json') as f:
12     linting = json.load(f)
13
14 # Define thresholds
15 THRESHOLDS = {
16     'max_complexity': 15,
17     'min_coverage': 70,
18     'max_linting_violations': 50,
19     'min_maintainability': 50
20 }
```



```

21
22 # Check thresholds
23 violations = []
24
25 # Check cyclomatic complexity
26 for module in complexity:
27     for func in module.get('functions', []):
28         if func['complexity'] > THRESHOLDS['max_complexity']:
29             violations.append(
30                 f"Function_{func['name']}_has_complexity_"
31                 f"{func['complexity']}_(max:{THRESHOLDS['max_complexity']})"
32             )
33
34 # Check test coverage
35 total_coverage = coverage['totals']['percent_covered']
36 if total_coverage < THRESHOLDS['min_coverage']:
37     violations.append(
38         f"Test_coverage_{total_coverage}%_is_below_"
39         f"minimum_{THRESHOLDS['min_coverage']}%"
40     )
41
42 # Check linting violations
43 violation_count = len(linting)
44 if violation_count > THRESHOLDS['max_linting_violations']:
45     violations.append(
46         f"Linting_violations_{violation_count}_exceed_"
47         f"maximum_{THRESHOLDS['max_linting_violations']}"
48     )

```

```

49
50 # Report violations
51 if violations:
52     print("Quality_gate_FAILED:")
53     for v in violations:
54         print(f"_{v}")
55     sys.exit(1)
56 else:
57     print("Quality_gate_PASSED")
58     sys.exit(0)

```

A.3 Planned Tool Configurations

A.3.1 Radon Configuration

Radon would analyze code complexity without additional configuration, but the following commands were planned:

Listing 3: Radon Analysis Commands

```

1 # Lines of Code analysis
2 radon raw . --json --output-file=radon_loc.json
3
4 # Cyclomatic Complexity
5 radon cc . --json --min=C --output-file=radon_complexity.json
6
7 # Maintainability Index
8 radon mi . --json --min=B --output-file=radon_maintainability.json
9
10 # Halstead metrics (code effort estimation)
11 radon hal . --json --output-file=radon_halstead.json

```

A.3.2 Pytest Configuration

Listing 4: pytest.ini

```
1 [pytest]
2 testpaths = tests
3 python_files = test_*.py
4 python_classes = Test*
5 python_functions = test_*
6
7 # Coverage settings
8 addopts =
9     --cov=src
10    --cov-report=html
11    --cov-report=json
12    --cov-report=term-missing
13    --cov-fail-under=70
14    --junitxml=junit.xml
15    --verbose
16
17 # Coverage exclusions
18 [coverage:run]
19 omit =
20     */tests/*
21     */venv/*
22     */migrations/*
23     */__pycache__/*
24
25 [coverage:report]
26 precision = 2
```

```
27 show_missing = True
28 skip_covered = False
```

A.3.3 Ruff Configuration

Listing 5: pyproject.toml - Ruff Configuration

```
1 [tool.ruff]
2 line-length = 100
3 target-version = "py311"
4
5 select = [
6     "E",    # pycodestyle errors
7     "W",    # pycodestyle warnings
8     "F",    # pyflakes
9     "I",    # isort
10    "C",    # flake8-comprehensions
11    "B",    # flake8-bugbear
12    "UP",   # pyupgrade
13    "N",    # pep8-naming
14 ]
15
16 ignore = [
17     "E501", # line too long (handled by formatter)
18     "B008", # function call in argument defaults
19 ]
20
21 [tool.ruff.per-file-ignores]
22 "__init__.py" = ["F401"] # unused imports
23 "tests/*.py" = ["S101"]  # allow assert statements
```

```
24
25 [tool.ruff.isort]
26 known-first-party = ["src"]
```

A.3.4 Mypy Configuration

Listing 6: mypy.ini

```
1 [mypy]
2 python_version = 3.11
3 warn_return_any = True
4 warn_unused_configs = True
5 disallow_untyped_defs = True
6 disallow_incomplete_defs = True
7 check_untyped_defs = True
8 no_implicit_optional = True
9 warn_redundant_casts = True
10 warn_unused_ignores = True
11 warn_no_return = True
12 strict_equality = True
13 show_error_codes = True
14
15 # Per-module options
16 [mypy-tests.*]
17 disallow_untyped_defs = False
18
19 [mypy-third_party_lib.*]
20 ignore_missing_imports = True
```

A.4 Planned Semantic Versioning Integration

The workflow would have used `python-semantic-release` for automated version management:

Listing 7: Semantic Release Workflow

```
1 name: Semantic Release
2
3 on:
4   push:
5     branches: [main]
6
7 jobs:
8   release:
9     runs-on: ubuntu-latest
10    steps:
11      - uses: actions/checkout@v3
12        with:
13          fetch-depth: 0
14
15      - name: Python Semantic Release
16        uses: relekang/python-semantic-release@master
17        with:
18          github_token: ${ secrets.GITHUB_TOKEN }
19          pypi_token: ${ secrets.PYPI_TOKEN }
```

Listing 8: pyproject.toml - Semantic Release Config

```
1 [tool.semantic_release]
2 version_variable = "src/__init__.py:__version__"
3 version_source = "tag"
```

```

4  commit_version_number = true
5  tag_commit = true
6  upload_to_pypi = false
7  upload_to_release = true
8
9  branch = "main"
10 hvcs = "github"
11
12 # Commit parsing
13 commit_parser = "angular"
14 major_on_zero = true
15 allow_zero_version = true
16
17 # Version bumping rules
18 [tool.semantic_release.commit_parser_options]
19 allowed_tags = [
20     "feat",      # New feature (minor bump)
21     "fix",       # Bug fix (patch bump)
22     "docs",      # Documentation only
23     "style",     # Code style changes
24     "refactor",  # Code refactoring
25     "perf",      # Performance improvements
26     "test",      # Test additions
27     "build",     # Build system changes
28     "ci",        # CI configuration
29 ]
30
31 major_tags = ["BREAKING CHANGE", "BREAKING"]
32 minor_tags = ["feat"]

```

```
patch_tags = ["fix", "perf"]
```

B Data Collection Scripts

The data collection process involved manually extracting data from GitHub's web interface and organizing it into CSV format. While no automated scraping scripts were used due to permission constraints, the following describes the manual collection workflow:

B.1 GitHub Web UI Navigation

For each repository (LangChain and Intel XPU Backend), data was collected through:

1. Pull Request Data Collection

- Navigate to repository → Pull Requests tab
- Filter by date range (e.g., October 29 - November 27, 2025 for LangChain)
- Filter by state: `is:pr is:open`, `is:pr is:merged`, `is:pr is:closed`
- Copy PR titles, numbers, and timestamps to CSV

2. Issue Data Collection

- Navigate to repository → Issues tab
- Filter by creation date within analysis window
- Identify bug-related issues using keyword search: `bug`, `error`, `fix`, `broken`, `crash`, `fail`, `regression`, `exception`
- Export issue titles and creation timestamps

3. Repository Metadata

- Navigate to Insights → Pulse for activity summary

- Record commit hashes and dates from commit history
- Note relative timestamps ("2 days ago") and convert to absolute dates

B.2 CSV Data Organization

Data was organized into four CSV files per repository:

Listing 9: CSV File Structure

```
1 # opened_requests.csv
2 msg,hash,date_rel,date_abs
3 "Add_authentication_feature",#12345,"2 days ago","2025-11-27"
4
5 # merged_requests.csv
6 msg,hash,date_rel,date_abs
7 "Fix_authentication_bug",#12346,"1 day ago","2025-11-28"
8
9 # closed_requests.csv
10 msg,hash,date_rel,date_abs
11 "Update_documentation",#12347,"3 hours ago","2025-11-29"
12
13 # opened_issues.csv
14 msg,hash,date_rel,date_abs
15 "Bug:_Login_fails_on_Safari",#678,"1 week ago","2025-11-22"
```

C Analysis Source Code

C.1 Key Functions

- `load_data(filepath)`: Loads CSV data and preprocesses timestamps

- Parses `date_abs` column into datetime objects
 - Handles missing or malformed timestamps
 - Returns pandas DataFrame
- `calculate_deployment_frequency(merged_df, days)`: Calculates DF metric
 - Counts merged PRs per day over analysis window
 - Computes mean, median, and weekly aggregate
 - Classifies performance: Elite ($>1/\text{day}$), High (weekly), Medium (monthly), Low ($<\text{monthly}$)
 - `calculate_lead_time(opened_df, merged_df)`: Calculates LT metric
 - Matches opened PRs with merged PRs using temporal proximity
 - Computes time delta in days
 - Returns median, mean, 25th/75th/90th percentiles
 - Classifies: Elite (<1 day), High (1-7 days), Medium (1-4 weeks), Low (>1 month)
 - `calculate_change_failure_rate(issues_df, merged_df)`: Calculates CFR metric
 - Identifies bug issues using keyword matching on issue titles/descriptions
 - Counts bug issues within analysis window
 - Computes ratio: $(\text{bug issues} / \text{total deployments}) \times 100\%$
 - Classifies: Elite (0-15%), High (16-30%), Medium (31-45%), Low ($>45\%$)
 - `calculate_time_to_restore(issues_df, merged_df)`: Calculates MTTR metric
 - Attempts to pair bug issues with fix PRs
 - Searches for issue references in PR titles/bodies
 - Computes time from bug opened to fix merged

- Returns median restoration time
- Classifies: Elite (<1 hour), High (<1 day), Medium (1-7 days), Low (>1 week)
- `generate_visualizations(metrics_dict, output_dir)`: Creates charts
 - Deployment frequency time series
 - Lead time distribution histogram
 - Change failure rate comparison (deployments vs bugs)
 - Performance classification summary

C.2 Usage Example

Listing 10: Running DORA Metrics Analysis

```

1  import pandas as pd
2  from dora_metrics import *
3
4  # Load data
5  opened = load_data('data/langchain/opened_requests.csv')
6  merged = load_data('data/langchain/merged_requests.csv')
7  closed = load_data('data/langchain/closed_requests.csv')
8  issues = load_data('data/langchain/opened_issues.csv')
9
10 # Calculate metrics
11 df = calculate_deployment_frequency(merged, days=30)
12 lt = calculate_lead_time(opened, merged)
13 cfr = calculate_change_failure_rate(issues, merged)
14 mttr = calculate_time_to_restore(issues, merged)
15
16 # Generate visualizations

```

```

17 metrics = {'df': df, 'lt': lt, 'cfr': cfr, 'mttr': mttr}
18 generate_visualizations(metrics, 'output/langchain/')
19
20 # Print summary
21 print(f"Deployment_Frequency:_{df['mean']:.2f}/day")
22 print(f"Lead_Time_(median):_{lt['median']:.2f}_days")
23 print(f"Change_Failure_Rate:_{cfr['rate']:.2f}%")
24 print(f"Time_to_Restore:_{mttr['median']:.2f}_hours")

```

D Raw Data Samples

D.1 LangChain Repository Data

Sample entries from the LangChain analysis dataset (October 29 - November 27, 2025):

Table 5: Sample Merged Pull Requests (LangChain)

PR #	Title	Date Relative	Date Absolute
#25843	anthropic[patch]: Release 0.41.1	2 days ago	2025-11-25
#25841	docs: Update Anthropic tools example	2 days ago	2025-11-25
#25839	core[patch]: Fix streaming in RunnableBinding	3 days ago	2025-11-24
#25832	community[patch]: Add timeout to Redis	4 days ago	2025-11-23
#25829	docs: Add citation format to papers	5 days ago	2025-11-22

Table 6: Sample Bug Issues (LangChain)

Issue #	Title	Date Relative	Date Absolute
#4521	Bug: Streaming fails with Azure OpenAI	1 day ago	2025-11-26
#4518	Error when using callback with agents	3 days ago	2025-11-24
#4510	Fix: Memory leak in conversation buffer	5 days ago	2025-11-22
#4502	Regression: Tool calling broken in 0.41.0	1 week ago	2025-11-20

D.2 Intel XPU Backend Data

Sample entries from Intel XPU Backend analysis dataset (November 4 - December 2, 2025):

Table 7: Sample Merged Pull Requests (Intel XPU Backend)

PR #	Title	Date Relative	Date Absolute
#2104	Support for Triton 3.2.0	1 day ago	2025-12-01
#2101	Fix GEMM kernel for FP16	2 days ago	2025-11-30
#2098	Add support for grouped convolutions	4 days ago	2025-11-28
#2095	Optimize reduction kernels	6 days ago	2025-11-26
#2092	Fix memory allocation in compiler	1 week ago	2025-11-25

E OpenTelemetry Configuration Files

E.1 Docker Compose Configuration

The Phase 2 OpenTelemetry approach used the following Docker Compose setup:

Listing 11: docker-compose.yml for OTel Stack

```

1 version: '3.8'
2
3 services:

```

```

4 otel-collector:
5     image: otel/opentelemetry-collector-contrib:latest
6     container_name: otel-collector
7     command: ["--config=/etc/otel-collector-config.yaml"]
8     volumes:
9         - ./otel-collector-config.yaml:/etc/otel-collector-config.yaml
10    ports:
11        - "4317:4317"    # OTLP gRPC receiver
12        - "4318:4318"    # OTLP HTTP receiver
13        - "8888:8888"    # Prometheus metrics
14    networks:
15        - otel-network
16    depends_on:
17        - openobserve
18
19 openobserve:
20     image: public.ecr.aws/zinclabs/openobserve:latest
21     container_name: openobserve
22     ports:
23         - "5080:5080"
24     environment:
25         - ZO_ROOT_USER_EMAIL=admin@example.com
26         - ZO_ROOT_USER_PASSWORD=admin123
27         - ZO_DATA_DIR=/data
28     volumes:
29         - openobserve-data:/data
30     networks:
31         - otel-network

```

```
33 networks:
34     otel-network:
35         driver: bridge
36
37 volumes:
38     openobserve-data:
```

E.2 OpenTelemetry Collector Configuration

Listing 12: otel-collector-config.yaml

```
1 receivers:
2     github:
3         github_org: langchain-ai
4         search_query: "org:langchain-ai is:pr"
5         scrapers:
6             github:
7                 collection_interval: 60s
8         auth:
9             authenticator: bearertokenauth
10
11 extensions:
12     bearertokenauth:
13         token: "${GITHUB_TOKEN}"
14
15 processors:
16     batch:
17         timeout: 10s
18         send_batch_size: 100
```

```

19
20 exporters:
21     otlphttp:
22         endpoint: "http://openobserve:5080/api/default/"
23         headers:
24             Authorization: "Basic_${OPENOBSERVE_CREDENTIALS}"
25             stream-name: "github-metrics"
26
27 service:
28     extensions: [bearertokenauth]
29     pipelines:
30         metrics:
31             receivers: [github]
32             processors: [batch]
33             exporters: [otlphttp]

```

E.3 Performance Classification Logic

Listing 13: DORA Classification

```

1 def classify_deployment_frequency(deployments_per_day):
2     """Elite: >1/day, High: weekly, Medium: monthly, Low: <monthly
3         """
4     if deployments_per_day >= 1:
5         return 'Elite'
6     elif deployments_per_day >= 1/7:
7         return 'High'
8     elif deployments_per_day >= 1/30:
9         return 'Medium'

```



```

9         else:
10             return 'Low'
11
12 def classify_lead_time(median_days):
13     """Elite: <1 day, High: 1-7 days, Medium: 1-4 weeks, Low: >1
14         month"""
15
16     if median_days < 1:
17         return 'Elite'
18
19     elif median_days <= 7:
20         return 'High'
21
22     elif median_days <= 28:
23         return 'Medium'
24
25     else:
26         return 'Low'
27
28
29 def classify_change_failure_rate(percentage):
30     """Elite: 0-15%, High: 16-30%, Medium: 31-45%, Low: >45%"""
31
32     if percentage <= 15:
33         return 'Elite'
34
35     elif percentage <= 30:
36         return 'High'
37
38     elif percentage <= 45:
39         return 'Medium'
40
41     else:
42         return 'Low'
43
44
45 def classify_time_to_restore(median_hours):
46     """Elite: <1 hour, High: <1 day, Medium: 1-7 days, Low: >1 week
47         """

```

```
36     if median_hours < 1:
37         return 'Elite'
38     elif median_hours < 24:
39         return 'High'
40     elif median_hours < 168: # 7 days
41         return 'Medium'
42     else:
43         return 'Low'
```